


The Expanding AI Attack Surface

ISC2 Detroit

Patrick Garrity

Security Researcher @ VulnCheck



AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies


AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies
Categorical Risk	Model Breakouts	Frontier Labs

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies
Categorical Risk	Model Breakouts	Frontier Labs

First Signs: Open Source (Github CVE Issuance)



Project Glasswing

Securing critical software
for the AI era


[Continue reading](#)

What Vulnerabilities Are Directly Attributable to Glasswing?

Despite the attention around Glasswing, only one publicly disclosed CVE can currently be directly tied to it. CVE-2026-4747 (FreeBSD NFS RCE) is explicitly attributed to Glasswing and Mythos Preview by name, described as fully autonomously identified and exploited.

The Glasswing page also references three vulnerabilities without CVE numbers: a 27-year-old OpenBSD flaw, a 16-year-old FFmpeg bug, and Linux kernel privilege escalation chains. All three are still under embargo pending patches.

Behind the Scenes Hardening Firefox with Claude Mythos Preview



By [Brian Grinstead](#), [Christian Holler](#), [Frederik Braun](#)


Posted on May 7, 2026 in [Featured Article](#), [Firefox](#), and [Security](#).

Two weeks ago we [announced](#) that we had identified and fixed an unprecedented number of latent security bugs in Firefox with the help of Claude Mythos Preview and other AI models. In this post, we'll go into more detail about how we approached this work, what we found, and advice for other projects on making good use of emerging capabilities to harden themselves against attack.

Suddenly, the bugs are very good

Firefox Security Bug Fixes by Month

All Sources • All Severities






CURL AND LIBCURL, SECURITY

MYTHOS FINDS A CURL VULNERABILITY

🕒 MAY 11, 2026 👤 DANIEL STENBERG 💬 33 COMMENTS

yes, as in singular *one*.


Cumulative CVE publications — top 20 CNAs by 2026 volume, 2023-2026
(y-axis scaled per panel; see legend for totals)



Check


**A FEW
MOMENTS LATER**

Cumulative CVE publications — top 30 CNAs by 2026 volume, 2023-2026
(y-axis scaled per panel; see legend for totals)



Check

Anthropic Ledger Launches May 22nd





Disclosure ledger

[Overview](#) · [Ledger](#) · [Archive](#) · [About](#)

[Light](#) [Dark](#) [System](#)

Once a finding has been validated by one of our external security research firms, a SHA-3-512 hash of the sealed report is published immediately as proof of possession. Status is shown once the maintainer has been notified; the identifier, project, and bug class are revealed only when the disclosure window closes. Committed entries that have not yet reached the maintainer show only the hash and commitment date.

MATCHING
1,611
of 1,611 entries

DISCLOSED
1,596

FIXED
27

PRE DISCLOSURE
15

WITHDRAWN
0

SEARCH

PROJECT

BUG CLASS

STATUS

hash, project, identifier, bug class...

All projects

All classes

All statuses

☐ Revealed only

Clear

DATE COMMITTED ▼	HASH / IDENTIFIER	PROJECT	BUG CLASS	STATUS
2026-05-20	c61fd338d6d28d96b01fd2b4...	—	—	DISCLOSED

Second Ledger Update

August 26th

(3-months later)

Disclosure ledger

[Overview](#) · [Ledger](#) · [Archive](#) · [About](#)

Light

Dark

System

Once a finding has been validated by one of our external security research firms, a SHA-3-512 hash of the sealed report is published immediately as proof of possession. Each ledger entry shows more detail as it progresses through disclosure: before the maintainer has been notified, only the commitment hash and the commitment date are published; once the report has been sent, the status and the severity assessments are shown and the discovery date is published (it drives the date filter); and the identifier, project, and bug class are revealed only when the disclosure window closes. A commitment that is withdrawn at any stage keeps its hash and commitment date and is marked withdrawn in the status column; a commitment withdrawn before the finding was revealed is also struck through in the table.

MATCHING

2,736

of 2,736 entries

DISCLOSED

2,300

FIXED

202

PRE DISCLOSURE

191

WITHDRAWN

245

MERGED

2

SEARCH

hash, project, identifier, bug class...

☐ Revealed only

Clear

PROJECT

All projects



BUG CLASS

All classes



STATUS

All statuses



ASSESSED BY

Any source



SEVERITY

Any severity



DISCOVERED

mm/dd/yyyy




to


mm/dd/yyyy




Anthropic Mythos ledger — 90-day disclosure deadline vs. patch released



Anthropic Mythos ledger — 90-day disclosure deadline vs. patch released




Anthropic Mythos ledger — cumulative findings over time




Anthropic Project Glasswing Findings

Anthropic Glasswing



What Percentage of Anthropic Discovered Vulnerabilities are Valid?


55% of the 449 revealed findings in Anthropic's Project Glasswing ledger have either been withdrawn or duplicate.




Revealed Findings as of August 28th, 2026

Anthropic Project Glasswing Severity Rating

Claude Severity



Maintainer Severity



ANTHROPIC CVE TRACKER


Overview

Tracking vulnerabilities that credit the Anthropic research team and are possibly discovered by [Project Glasswing](#).

CURRENT CVE COUNT: 225

Distributions

CVSS Severity



225 CVEs

Exploited in the Wild

Source: VulnCheck KEV



225 CVEs

Open AI Patch The Planet



LAST UPDATED · SEP 11, 2026

The State of Patch the Planet

Bugs found, patches shipped, and fixes merged at the speed of AI across all the open-source projects in the initiative.

Apply now



About the initiative ↗

ISSUES FOUND

1,829

POTENTIAL BUGS REPORTED

AWAITING PATCH

1,211

REPORTED + CONFIRMED

PATCHES OPEN

203

FIX OPEN UPSTREAM

PATCHES ACCEPTED

415


MERGED UPSTREAM

PROJECTS

60


CODEBASES UNDER REVIEW

Issues over time




PEAK: 224 ON JUN 2 · INCLUDES ALL REPORTS FILED, INCLUDING REPORTS CLOSED DURING TRIAGE

Report pipeline



Patch activity

■ PATCH THE PLANET ★ PTP BEGINS ■ BOOKED



INCLUDES ALL REPORTS FILED, INCLUDING REPORTS CLOSED DURING TRIAGE

Accepted fixes by severity

415 TOTAL



■ HIGH **62** 14.9% ■ MEDIUM **196** 47.2% ■ LOW **85** 20.5% ■ INFO **31** 7.5% ■ UNDETERMINED **41** 9.9%

Vulnerability Clearing Warehouses

THE **LINUX** FOUNDATION PROJECTS



RELEASES

White House Launches Gold Eagle Initiative for Unprecedented Cybersecurity Vulnerability Coordination

The White House | July 14, 2026

ATHENA



America's Cyber Defense Agency

NATIONAL COORDINATOR FOR CRITICAL INFRASTRUCTURE SECURITY AND RESILIENCE





r/TechNadu · 1d ago

technadu

Human




Top 1% Poster



65% of surveyed bug hunters say they've withheld a vulnerability because there was no clear reporting path

Vulnerability Discovery Evidence (CVE's Published)




AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies
Categorical Risk	Model Breakouts	Frontier Labs

Known Exploited Vulnerabilities

Time From CVE to Exploitation Evidence


Source: VulnCheck KEV 2026



Are AI Discovered Vulnerabilities More Dangerous?

What Percentage of Vulnerabilities Discovered using AI Are Exploited?

1.3% of the 1,734 CVEs
Identified by Berkley
Vulnerability Initiative and
Patrick Garrity's Anthropic
Credited CVEs project are known
to be exploited.



What products were exploited using AI discovered vulnerabilities?


CVEs Identified by Berkley Vulnerability Initiative and Patrick Garrity's Anthropic Credited CVEs project are known to be exploited.

Microsoft Windows	Linux Kernel	BeyondTrust	openclaw
proftpd		GeoServer	
Apache ActiveMQ	Slider Future	chef	chancms
ghost			

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies
Categorical Risk	Model Breakouts	Frontier Labs

Threat Actor Harnesses AI Models for Autonomous Cyberattacks



LLM Attack Reasoning...

All three Langflow need public flow ID but no auto_login — stuck. Deployments small (84 alive), exploitable probably 0. Search for larger-scale vulns.

The n8n one with 258 stars and CVSS 10.0 looks extremely promising! n8n is a very popular workflow automation tool similar to Langflow, and it has way more stars on the PoC which suggests it's widely exploited.


Actually, looking at the exploit code again — it needs an UNAUTHENTICATED form with file upload. These forms have auth enabled. Let me check other targets.

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies
Categorical Risk	Model Breakouts	Frontier Labs


Top Targeted Technologies

Source: VulnCheck KEV



Top Targeted Technologies

Source: VulnCheck KEV
(1H-2026)



AI Products Exploited In 1H-2026

AI Product							
Langflow				Praison AI	openclaw	Mintplexlabs anythin...	
				Ollama	n8n	MCPJam Inspector	
						lfprojects mlflow	
Majordomo		lollms		internlm lm...	Gradio	comfy	budibase
BerriAI LiteLLM		dify		FlowiseAI		anyscale ray	

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Evidence
Vulnerability Discovery	Easier to discovery Vulnerabilities	100% YoY Increase in CVEs
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Exploitation Timeline Dropping 120 > 80 Days
Attack Orchestration	Automated Campaign Activity	UAT-10147 / GT-1002 (Chinese Campaigns)
Attack Surface	New Exposure / Priveledged Access	AI Top 5 Attacked Technologies
Categorical Risk	Model Breakouts	Frontier Labs

Frontier Model Breakout Attacks

July 21, 2026 Security

OpenAI and Hugging Face
partner to address security
incident during model evaluation

ANTHROPIC

In a review of our cybersecurity evaluation transcripts, we found three incidents in which a Claude model reached the internet from within or while interacting with a third-party evaluation environment, and then gained unauthorized access to the real systems of three different organizations.

 Meta

**Meta AI breaches external firm during
security testing sandbox error**

AUGUST 8, 2026 · 7:43 AM ET

 VulnCheck

What Actions Can Defenders Take?

AI Vulnerability & Exploitation Risk Considerations

Risk Category	AI's Impact	Actions
Vulnerability Discovery	Easier to discovery Vulnerabilities	Expand Vulnerability Detection & Remediation. Use AI Harness / AI Models to expand visibility into undiscovered vulnerabilities.
Vulnerability Exploitation	Faster Exploit Development / Weaponization	Accelerate Remediation / Mitigating Controls / Root Cause Analysis
Attack Orchestration	Automated Campaign Activity	Detection & Response
Attack Surface	New Exposure / Priveledged Access	Zero Trust / Least Privledged Access / Procurement Consideration
Categorical Risk	Model Breakouts	Audit for Misconfigurations

Questions?



Patrick Garrity

SECURITY RESEARCHER

Follow My Work

